

---

# 5D Neural Network Interpolator Documentation

*Release 0.1.0*

Makimona Kiakisolako (bamk3)

Nov 26, 2025



# User Guide

<b>1 Overview</b>	<b>3</b>
<b>2 Key Features</b>	<b>5</b>
2.1 Application Features . . . . .	5
2.2 Development Features . . . . .	5
<b>3 Quick Start</b>	<b>7</b>
<b>4 Documentation Contents</b>	<b>9</b>
4.1 Installation Guide . . . . .	9
4.2 Quick Start Guide . . . . .	15
4.3 Usage Guide . . . . .	20
4.4 Dataset Specifications . . . . .	27
4.5 Backend API Reference . . . . .	34
4.6 Frontend Components . . . . .	41
4.7 Neural Network Module . . . . .	50
4.8 Performance and Profiling . . . . .	54
4.9 Testing Overview . . . . .	64
4.10 Local Deployment Guide . . . . .	71
4.11 System Architecture . . . . .	77
<b>5 Indices and Tables</b>	<b>89</b>
<b>6 Project Information</b>	<b>91</b>
6.1 Links . . . . .	91
<b>Python Module Index</b>	<b>93</b>



Welcome to the **5D Neural Network Interpolator** documentation. This application provides a complete solution for 5D function interpolation using neural networks, developed as coursework for the DIS course at the University of Cambridge.

1 2

---

<sup>1</sup> <https://www.python.org/downloads/release/python-312/>  
<sup>2</sup> <https://nextjs.org/>



# 1

## Overview

The 5D Interpolator is a full-stack web application that enables:

- **Fast Neural Network Training:** CPU-optimized training in under 1 minute on datasets up to 10,000 samples
- **Configurable Architecture:** Fully customizable hyperparameters including layer sizes, learning rate, and iterations
- **Interactive Interface:** Modern React-based UI with real-time feedback
- **Batch & Single Predictions:** Support for both bulk dataset predictions and individual feature inputs
- **RESTful API:** Complete FastAPI backend with automatic documentation



# 2

## Key Features

### 2.1 Application Features

- 5D neural network interpolation with configurable architecture
- Dataset upload (.pkl format) with automatic validation
- Model training with customizable hyperparameters via sliders
- Single and batch predictions
- RESTful API with OpenAPI/Swagger documentation
- Modern, responsive UI with dark mode support

### 2.2 Development Features

- Docker containerization with hot reload
- Comprehensive test suite (52 tests, 74% coverage)
- Multi-stage Docker builds for development and production
- Environment-based configuration
- Helper scripts for common operations



## Quick Start

### Using Docker (Recommended)

```
# Complete setup from scratch
./scripts/docker-start.sh

# Access the application
# Frontend: http://localhost:3000
# Backend API: http://localhost:8000
# API Documentation: http://localhost:8000/docs
```

### Manual Setup

```
# Backend
cd backend
pip install -r requirements.txt
unicorn main:app --reload

# Frontend (separate terminal)
cd frontend
npm install
npm run dev
```

### Download Documentation

This documentation is available in multiple formats:

```
# Generate PDF and HTML archives
./scripts/build-docs-pdf.sh
```

Available formats:

- **PDF (~419 KB)** - For offline reading and printing
- **HTML Archive (~8.2 MB)** - Complete offline browsable documentation
- **Online HTML** - This current format

See *Installation Guide* for details on building and downloading documentation.



## Documentation Contents

### 4.1 Installation Guide

This guide covers all methods for installing and running the 5D Neural Network Interpolator.

#### 4.1.1 Prerequisites

##### System Requirements

- **Operating System:** macOS, Linux, or Windows (with WSL2)
- **RAM:** Minimum 4GB (8GB recommended)
- **Disk Space:** Minimum 2GB free space
- **Internet Connection:** Required for initial setup

##### Software Requirements

###### For Docker Installation (Recommended):

- Docker 20.10+ or Docker Desktop
- Docker Compose v2.0+

###### For Manual Installation:

- Python 3.12+
- Node.js 20+
- npm 10.8+
- pip 23+

#### 4.1.2 Installation Methods

##### Method 1: Docker Installation (Recommended)

This is the fastest and most reliable method.

###### Step 1: Verify Docker is Running

```
# Check Docker is installed and running
docker --version
docker compose version

# On Linux, ensure Docker service is running
sudo systemctl status docker
```

### Step 2: Clone the Repository

```
git clone <repository-url>
cd interpolator
```

### Step 3: Run Setup Script

```
# Complete setup (clean + rebuild + start)
./scripts/docker-start.sh
```

This script will:

1. Check if Docker is running
2. Clean up any existing containers
3. Create environment configuration
4. Build Docker images (~3-5 minutes)
5. Start all services
6. Display access URLs

### Step 4: Verify Installation

```
# Check service status
docker compose ps

# Test backend health
curl http://localhost:8000/health

# Test frontend (should return HTML)
curl http://localhost:3000
```

**Access URLs:**

- Frontend: <http://localhost:3000>
- Backend API: <http://localhost:8000>
- API Documentation: <http://localhost:8000/docs>

## Method 2: Manual Installation

### Step 1: Install Backend Dependencies

```
cd backend

# Create virtual environment (recommended)
python3 -m venv venv
source venv/bin/activate # On Windows: venv\Scripts\activate
```

(continues on next page)

(continued from previous page)

```
# Install dependencies
pip install -r requirements.txt
```

### Step 2: Install Frontend Dependencies

```
cd frontend
npm install
```

### Step 3: Start Backend Server

```
cd backend
source venv/bin/activate # If using virtual environment
uvicorn main:app --reload 0.0.0.0 --port 8000
```

### Step 4: Start Frontend Server (in new terminal)

```
cd frontend
npm run dev
```

#### Access URLs:

- Frontend: <http://localhost:3000>
- Backend API: <http://localhost:8000>
- API Documentation: <http://localhost:8000/docs>

## 4.1.3 Environment Configuration

### Environment Variables

The application uses environment variables for configuration. Three preset files are provided:

- .env.development - For local development
- .env.production - For production deployment
- .env.example - Template with all available variables

### Key Variables:

```
# Backend
BACKEND_PORT=8000
CORS_ORIGINS=http://localhost:3000

# Frontend
FRONTEND_PORT=3000
NEXT_PUBLIC_API_URL=http://localhost:8000

# Docker
BUILD_TARGET=development # or 'production'
```

### Setup:

```
# Copy development configuration  
cp .env.development .env  
  
# Or for production  
cp .env.production .env
```

#### 4.1.4 Troubleshooting

##### Docker Issues

“Docker is not running” error:

```
# macOS  
open -a Docker  
  
# Linux  
sudo systemctl start docker  
  
# Check status  
docker info
```

“Port already in use” error:

```
# Find process using port 3000  
lsof -i :3000  
  
# Kill the process  
kill -9 <PID>
```

“docker-compose: command not found”:

You have Docker Compose v2 (plugin version). Use:

```
docker compose # (with space, not hyphen)
```

##### Permission Issues (Linux)

```
# Add user to docker group  
sudo usermod -aG docker $USER  
  
# Apply changes  
newgrp docker  
  
# Verify  
docker ps
```

##### Python/Node Issues

Wrong Python version:

```
# Check version  
python3 --version
```

(continues on next page)

(continued from previous page)

```
# Install Python 3.12 via package manager
# macOS:
brew install python@3.12

# Ubuntu/Debian:
sudo apt install python3.12
```

**npm install fails:**

```
# Clear npm cache
npm cache clean --force

# Delete node_modules and retry
rm -rf node_modules package-lock.json
npm install
```

#### 4.1.5 Verifying Installation

Run the complete test suite to verify everything works:

```
# Using Docker
./scripts/docker-dev.sh test-backend

# Manual installation
cd backend
pytest
```

Expected output: 52 passed with 74% coverage

#### 4.1.6 Building Documentation

This documentation can be built locally for offline access:

##### Quick Build

```
./scripts/build-docs.sh
```

This automated script will:

1. Check Python installation (3.12+ required)
2. Create virtual environment for Sphinx
3. Install Sphinx and dependencies
4. Build HTML documentation
5. Open in your default browser

The documentation will be available at:

```
docs/build/html/index.html
```

### Manual Build

For manual control over the build process:

```
cd docs

# Create virtual environment (first time only)
python3 -m venv venv
source venv/bin/activate

# Install dependencies (first time only)
pip install -r requirements.txt

# Build documentation
sphinx-build -b html source build/html

# Open in browser
open build/html/index.html # macOS
xdg-open build/html/index.html # Linux
```

### Rebuilding Documentation

To rebuild after making changes:

```
cd docs
source venv/bin/activate

# Clean previous build
rm -rf build/html

# Rebuild
sphinx-build -b html source build/html
```

### Documentation Requirements

The documentation build requires:

- Python 3.12+
- Sphinx 8.2.3+
- sphinx-rtd-theme
- sphinxcontrib packages

All dependencies are listed in `docs/requirements.txt`

### Downloading Documentation

**Generate Downloadable Documentation:**

```
./scripts/build-docs-pdf.sh
```

This generates multiple downloadable formats:

- **PDF:** `docs/build/downloads/5D-Interpolator-Documentation.pdf` (~419 KB)
- **HTML Archive (tar.gz):** `docs/build/downloads/5D-Interpolator-Documentation-HTML.tar.gz` (~8.2 MB)

- **HTML Archive (zip):** docs/build/downloads/5D-Interpolator-Documentation-HTML.zip (~8.2 MB)

#### PDF Generation Requirements:

For LaTeX-based PDF (recommended):

- **macOS:** Install MacTeX

```
brew install --cask mactex
```

- **Ubuntu/Debian:**

```
sudo apt-get install texlive-latex-extra texlive-fonts-recommended
```

- **Fallback:** If LaTeX not available, script automatically uses rst2pdf

#### Using Downloaded Documentation:

- **PDF:** Open directly in any PDF reader
- **HTML Archives:** Extract and open index.html in a web browser

```
# Extract tar.gz
tar -xzf 5D-Interpolator-Documentation-HTML.tar.gz
open html/index.html

# Or extract zip
unzip 5D-Interpolator-Documentation-HTML.zip
open index.html
```

### 4.1.7 Next Steps

- *Quick Start Guide* - Get started with your first model
- *Usage Guide* - Learn about features and workflows
- *Dataset Specifications* - Understand dataset requirements

## 4.2 Quick Start Guide

Get started with the 5D Neural Network Interpolator in just a few minutes.

### 4.2.1 Overview

This guide walks you through:

1. Starting the application
2. Uploading a training dataset
3. Training a model with custom hyperparameters
4. Making predictions

### 4.2.2 Starting the Application

Using Docker (recommended)

```
# Start all services  
./scripts/docker-start.sh  
  
# Or for quick start (if already set up)  
./scripts/docker-dev.sh  
  
# To stop services  
./scripts/docker-stop.sh
```

### Soft Manual Start (Using local-build shell script)

```
# Start the backend and frontend using local-build script  
./scripts/local-build.sh  
  
# To stop services  
./scripts/local-stop.sh
```

## Manual Start

### Terminal 1 - Backend:

```
cd backend  
unicorn main:app --reload
```

### Terminal 2 - Frontend:

```
cd frontend  
npm run dev
```

## Access the Application

Open your browser and navigate to:

- **Frontend:** <http://localhost:3000>
- **API Docs:** <http://localhost:8000/docs>

## Build the documentation

```
./scripts/build-docs.sh
```

### 4.2.3 Step-by-Step Workflow

#### Step 1: Upload Training Dataset

1. Navigate to <http://localhost:3000/upload>
2. Select “Training” dataset type
3. Click “Click to upload” or drag and drop your .pk1 file
4. Wait for validation and preview
5. Click “Proceed to Training →”

#### Dataset Requirements:

- Format: Python pickle (.pkl)
- Structure: Dictionary with keys 'X' and 'y'
- X: NumPy array of shape (n\_samples, 5) - 5D feature vectors
- y: NumPy array of shape (n\_samples,) - 1D target values

#### Example Dataset Creation:

```
import numpy as np
import pickle

# Generate sample data
n_samples = 1000
X = np.random.randn(n_samples, 5)
y = np.sum(X**2, axis=1) + 0.1 * np.random.randn(n_samples)

# Save as pickle
data = {'X': X, 'y': y}
with open('training_data.pkl', 'wb') as f:
    pickle.dump(data, f)
```

#### Step 2: Configure Hyperparameters

On the training page, you'll see interactive sliders for:

##### Neural Network Architecture:

- **Hidden Layer 1:** 8-256 neurons (default: 64)
- **Hidden Layer 2:** 8-128 neurons (default: 32)
- **Hidden Layer 3:** 4-64 neurons (default: 16)

##### Training Parameters:

- **Learning Rate:** 0.0001-0.01 (default: 0.001)
- **Max Iterations:** 100-2000 (default: 500)
- **Early Stopping:** On/Off (default: On)

##### Tips:

- Larger networks (more neurons) = more capacity but slower training
- Higher learning rates = faster convergence but may be unstable
- Early stopping prevents overfitting and saves time

#### Step 3: Train the Model

1. Adjust hyperparameters using the sliders
2. Click “**Start Training**”
3. Wait for training to complete (<1 minute for typical datasets)
4. Review the results:
  - **R<sup>2</sup> Score:** Model fit quality (>0.95 is excellent)
  - **MSE/MAE/RMSE:** Error metrics

- **Hyperparameters Used:** Confirmation of settings

#### Training Results Example:

```
Training Complete
```

```
R2 Score: 0.9876
MSE:      0.0123
MAE:      0.0987
RMSE:     0.1109
```

#### Hyperparameters Used:

```
Architecture: [64, 32, 16]
Learning Rate: 0.001
Max Iterations: 500
Early Stopping: Yes
```

### Step 4: Make Predictions

#### Option A: Batch Prediction

1. Navigate to <http://localhost:3000/upload>
2. Select “Prediction” dataset type
3. Upload a .pkl file containing only X data (shape: n, 5)
4. Go to <http://localhost:3000/predict>
5. Select “Batch Prediction” mode
6. Click “Generate Batch Predictions”

#### Example Prediction Dataset:

```
import numpy as np
import pickle

# Generate prediction inputs
X_pred = np.random.randn(100, 5)

# Save as pickle
with open('prediction_data.pkl', 'wb') as f:
    pickle.dump(X_pred, f)
```

#### Option B: Single Prediction

1. Go to <http://localhost:3000/predict>
2. Select “Single Prediction” mode
3. Enter values for all 5 features
4. Click “Predict”
5. View the result

#### Example Single Prediction:

**Input Features:**

F1: 1.2345  
 F2: -0.5678  
 F3: 0.9876  
 F4: -1.2345  
 F5: 0.5432

**Prediction Result:**

3.456789

#### 4.2.4 Common Workflows

##### Experiment with Hyperparameters

1. Upload dataset
2. Train with default settings
3. Note  $R^2$  score
4. Upload SAME dataset again (resets training state)
5. Adjust hyperparameters
6. Train again
7. Compare results

##### Quick Iteration Cycle

```
# Using Docker - complete reset
./scripts/docker-start.sh

# Or just restart services
./scripts/docker-dev.sh restart
```

#### 4.2.5 Best Practices

##### Dataset Preparation

- **Size:** 1,000-10,000 samples recommended
- **Quality:** Remove NaN/inf values before upload
- **Normalization:** Not required (automatic standardization)
- **Validation:** Check data shape and types before upload

##### Model Training

- Start with default hyperparameters
- Adjust based on  $R^2$  score:
  - $R^2 < 0.8$ : Increase network size or iterations
  - $R^2 > 0.99$ : May be overfitting, reduce complexity
  - Training too slow: Reduce iterations or network size
- Early stopping is recommended for most cases

### Performance Tips

- Use Docker for consistent performance
- Train on datasets < 10,000 samples for <1min training
- Batch predictions are faster than many single predictions
- Keep browser tab active during training

## 4.2.6 Troubleshooting Quick Start Issues

“Upload Dataset First” button stuck:

- Refresh the page
- Check backend is running: <http://localhost:8000/health>
- Re-upload the dataset

Training fails:

- Verify dataset format (must be dictionary with ‘X’ and ‘y’)
- Check dataset shape (X must be  $n \times 5$ , y must be 1D)
- Try with smaller dataset first

Predictions fail:

- Ensure model is trained first
- For batch: upload prediction dataset
- For single: fill all 5 feature fields

## 4.2.7 Next Steps

- [Usage Guide](#) - Detailed feature documentation
- [Dataset Specifications](#) - Dataset format specifications
- [Backend API Reference](#) - API reference for programmatic access
- [Testing Overview](#) - Running tests

## 4.3 Usage Guide

Comprehensive guide to using the 5D Neural Network Interpolator.

### 4.3.1 Application Workflow

The typical workflow consists of three main steps:

1. **Upload Training Dataset** → 2. **Train Model** → 3. **Make Predictions**

### 4.3.2 Step 1: Upload Training Dataset

Navigate to the Upload page and select a training dataset.

## Dataset Requirements

The training dataset must be a Python pickle file (.pkl) containing:

```
{
    'X': numpy.ndarray, # Shape: (n_samples, 5)
    'y': numpy.ndarray # Shape: (n_samples,)
}
```

Where:

- X: 5-dimensional feature vectors (independent variables)
- y: 1-dimensional target values (dependent variable)
- n\_samples: Number of training examples

## Example Dataset Creation

```
import numpy as np
import pickle

# Generate 1000 samples
n_samples = 1000

# Create 5D features
X = np.random.randn(n_samples, 5)

# Create target (example: sum of squares)
y = np.sum(X**2, axis=1) + 0.1 * np.random.randn(n_samples)

# Save as pickle
data = {'X': X, 'y': y}
with open('training_data.pkl', 'wb') as f:
    pickle.dump(data, f)
```

## Upload Process

1. Click “Training” dataset type
2. Click upload area or drag file
3. Wait for validation
4. Review data preview showing:
  - Total samples
  - Data shape
  - First 5 rows
5. Click “Proceed to Training →”

### 4.3.3 Step 2: Train Model

Configure hyperparameters and train the neural network.

## Hyperparameter Configuration

### Neural Network Architecture

- **Hidden Layer 1** (8-256 neurons)
  - Controls first layer capacity
  - Default: 64 neurons
  - Larger = more complex patterns
- **Hidden Layer 2** (8-128 neurons)
  - Controls second layer capacity
  - Default: 32 neurons
  - Typically smaller than layer 1
- **Hidden Layer 3** (4-64 neurons)
  - Controls third layer capacity
  - Default: 16 neurons
  - Smallest layer before output

### Training Parameters

- **Learning Rate** (0.0001-0.01)
  - Speed of gradient descent
  - Default: 0.001
  - Higher = faster but less stable
  - Lower = slower but more precise
- **Max Iterations** (100-2000)
  - Maximum training epochs
  - Default: 500
  - Higher = more training time
  - May stop early if enabled
- **Early Stopping** (On/Off)
  - Stops when validation loss plateaus
  - Default: On (recommended)
  - Prevents overfitting
  - Saves computation time

## Recommended Configurations

### Default (Balanced)

```
Architecture: [64, 32, 16]
Learning Rate: 0.001
Max Iterations: 500
Early Stopping: On
```

(continues on next page)

(continued from previous page)

Use for: Most datasets  
Expected time: 15-30 seconds

### Fast Training

Architecture: [32, 16, 8]  
Learning Rate: 0.01  
Max Iterations: 200  
Early Stopping: On

Use for: Quick experiments  
Expected time: 5-10 seconds

### High Accuracy

Architecture: [128, 64, 32]  
Learning Rate: 0.001  
Max Iterations: 1000  
Early Stopping: On

Use for: Best possible fit  
Expected time: 30-60 seconds

### Training Process

1. Adjust sliders to desired values
2. Click “**Start Training**”
3. Wait for training (typically <1 minute)
4. Review results

### Understanding Training Results

After training completes, you'll see:

#### Performance Metrics

- **R<sup>2</sup> Score:** Model fit quality (0-1)
  - >0.95: Excellent
  - 0.90-0.95: Very good
  - 0.80-0.90: Good
  - <0.80: May need tuning
- **MSE (Mean Squared Error):** Average squared error
  - Lower is better
  - Scale depends on target values
- **MAE (Mean Absolute Error):** Average absolute error
  - Lower is better
  - Same units as target variable

- **RMSE (Root Mean Squared Error):** Square root of MSE
  - Lower is better
  - Same units as target variable

### Hyperparameters Used

Displays the configuration used for training.

### 4.3.4 Step 3: Make Predictions

Two prediction modes are available.

#### Batch Prediction

For predicting multiple samples at once.

##### Dataset Requirements:

```
# Pickle file containing only X data
X_pred = numpy.ndarray # Shape: (n_samples, 5)
```

##### Example:

```
import numpy as np
import pickle

# Create prediction inputs
X_pred = np.random.randn(100, 5)

# Save as pickle
with open('prediction_data.pkl', 'wb') as f:
    pickle.dump(X_pred, f)
```

##### Steps:

1. Upload prediction dataset (Upload page)
2. Go to Predict page
3. Select “Batch Prediction”
4. Click “Generate Batch Predictions”
5. View results

#### Single Prediction

For predicting one sample at a time.

##### Steps:

1. Go to Predict page
2. Select “Single Prediction”
3. Enter values for all 5 features
4. Click “Predict”
5. View result

##### Example Input:

```
Feature 1: 1.2345
Feature 2: -0.5678
Feature 3: 0.9876
Feature 4: -1.2345
Feature 5: 0.5432
```

Result: 3.456789

### 4.3.5 Advanced Usage

#### Experimenting with Hyperparameters

To compare different configurations:

1. Train with configuration A
2. Note R<sup>2</sup> score
3. Go to Upload page
4. Re-upload SAME dataset (resets state)
5. Return to Train page
6. Train with configuration B
7. Compare results

#### API Usage

For programmatic access, use the REST API:

```
import requests

BASE_URL = "http://localhost:8000"

# Upload dataset
with open('data.pkl', 'rb') as f:
    r = requests.post(
        f"{BASE_URL}/upload-fit-dataset/",
        files={'file': f}
    )

# Train with custom hyperparameters
r = requests.post(
    f"{BASE_URL}/start-training/",
    json={
        "hyperparameters": {
            "hidden_layer_1": 128,
            "learning_rate": 0.01
        }
    }
)

# Make prediction
r = requests.post(
    f"{BASE_URL}/predict-single/",
```

(continues on next page)

(continued from previous page)

```
    json={"features": [1, 2, 3, 4, 5]}\n}\n\nprint(r.json())
```

See [Backend API Reference](#) for complete API reference.

### 4.3.6 Tips and Best Practices

#### Dataset Preparation

- Remove NaN/inf values before upload
- Ensure consistent data types
- Check for outliers
- Recommended size: 1,000-10,000 samples

#### Model Training

- Start with defaults
- Increase complexity if  $R^2 < 0.9$
- Reduce complexity if overfitting
- Use early stopping for efficiency
- Monitor training time

#### Prediction

- Ensure prediction data matches training scale
- Use batch mode for efficiency
- Single mode good for testing
- Validate results against known values

### 4.3.7 Troubleshooting

#### Training Issues

##### $R^2$ score too low (<0.8):

- Increase network size
- Increase iterations
- Try different learning rate
- Check data quality

##### Training too slow (>60 seconds):

- Reduce network size
- Reduce iterations
- Enable early stopping
- Use smaller dataset

**Model fails to converge:**

- Reduce learning rate
- Increase iterations
- Check for data issues

**Prediction Issues****Predictions seem wrong:**

- Verify model is trained
- Check prediction data format
- Ensure feature scales match training
- Review R<sup>2</sup> score

**Batch prediction fails:**

- Verify data shape (n, 5)
- Check file format (.pkl)
- Ensure model is trained

### 4.3.8 Keyboard Shortcuts

While using the application:

- **Refresh page:** Reset state
- **Browser back:** Navigate between pages
- **Ctrl/Cmd + Click link:** Open in new tab

### 4.3.9 Next Steps

- *Backend API Reference* - API reference
- *Testing Overview* - Run tests
- *Dataset Specifications* - Dataset specifications

## 4.4 Dataset Specifications

Complete specification for dataset formats used by the 5D Interpolator.

### 4.4.1 Training Dataset Format

**Structure**

Training datasets must be Python pickle files containing a dictionary:

```
{
    'X': numpy.ndarray, # Feature matrix
    'y': numpy.ndarray # Target vector
}
```

## Requirements

### File Format:

- Extension: .pkl
- Type: Python pickle file
- Encoding: Binary

### X (Features):

- Type: numpy.ndarray
- Shape: (n\_samples, 5)
- Dtype: float32 or float64
- Values: Any real numbers (will be standardized)
- Constraints:
  - Must have exactly 5 features
  - No NaN or inf values
  - At least 100 samples recommended

### y (Targets):

- Type: numpy.ndarray
- Shape: (n\_samples,) - 1D array
- Dtype: float32 or float64
- Values: Any real numbers
- Constraints:
  - Must match number of samples in X
  - No NaN or inf values

## Example Creation

```
import numpy as np
import pickle

# Generate features (1000 samples, 5 features)
X = np.random.randn(1000, 5)

# Generate targets
y = np.sum(X**2, axis=1)

# Create dataset dictionary
dataset = {'X': X, 'y': y}

# Save as pickle
with open('training_data.pkl', 'wb') as f:
    pickle.dump(dataset, f)
```

## Validation

The system automatically validates:

- ✓ File is readable pickle
- ✓ Contains 'X' and 'y' keys
- ✓ X has shape (n, 5)
- ✓ y has shape (n,)
- ✓ X and y have same number of samples
- ✓ No NaN or inf values

## 4.4.2 Prediction Dataset Format

### Structure

Prediction datasets must be Python pickle files containing a NumPy array:

```
numpy.ndarray # Shape: (n_samples, 5)
```

### Requirements

#### File Format:

- Extension: .pkl
- Type: Python pickle file
- Encoding: Binary

#### Data:

- Type: numpy.ndarray
- Shape: (n\_samples, 5)
- Dtype: float32 or float64
- Values: Any real numbers
- Constraints:
  - Must have exactly 5 features
  - No NaN or inf values
  - Any number of samples

### Example Creation

```
import numpy as np
import pickle

# Generate prediction inputs (100 samples, 5 features)
X_pred = np.random.randn(100, 5)

# Save as pickle
with open('prediction_data.pkl', 'wb') as f:
    pickle.dump(X_pred, f)
```

## 4.4.3 Data Preprocessing

### Automatic Standardization

The system automatically standardizes all features using:

```
from sklearn.preprocessing import StandardScaler  
  
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```

This means:

- Each feature is centered (mean = 0)
- Each feature is scaled (std = 1)
- Same transformation applied to predictions
- No manual normalization needed

## Data Splitting

Training data is automatically split:

- 60% Training set
- 20% Validation set
- 20% Test set

Split is random with fixed seed (42) for reproducibility.

### 4.4.4 Best Practices

#### Dataset Size

##### Recommended Sizes:

- Minimum: 100 samples
- Optimal: 1,000-10,000 samples
- Maximum: No hard limit (training time increases)

##### Training Time by Size:

- 100 samples: ~5 seconds
- 1,000 samples: ~15 seconds
- 10,000 samples: ~45 seconds
- 100,000 samples: ~5 minutes

#### Data Quality

##### Check for:

- Missing values (NaN)
- Infinite values (inf)
- Outliers (>3 std from mean)
- Data type consistency
- Correct dimensions

##### Example Validation:

```

import numpy as np

def validate_dataset(X, y):
    """Validate dataset before saving"""

    # Check shapes
    assert X.ndim == 2, "X must be 2D"
    assert X.shape[1] == 5, "X must have 5 features"
    assert y.ndim == 1, "y must be 1D"
    assert X.shape[0] == y.shape[0], "X and y must have same samples"

    # Check for invalid values
    assert not np.any(np.isnan(X)), "X contains NaN"
    assert not np.any(np.isinf(X)), "X contains inf"
    assert not np.any(np.isnan(y)), "y contains NaN"
    assert not np.any(np.isinf(y)), "y contains inf"

    print(f"✓ Dataset valid: {X.shape[0]} samples, 5 features")

# Use it
validate_dataset(X, y)

```

## Feature Engineering

Consider:

- Feature scaling (optional, auto-standardized)
- Polynomial features for non-linear relationships
- Interaction terms
- Domain-specific transformations

### 4.4.5 Common Use Cases

#### Regression Problems

##### Example: Function Approximation

```

# Approximate function: f(x1,...,x5) = x1^2 + x2*x3 - x4 + sin(x5)
import numpy as np

n = 1000
X = np.random.randn(n, 5)
y = X[:, 0]**2 + X[:, 1]*X[:, 2] - X[:, 3] + np.sin(X[:, 4])

```

#### Scientific Data

##### Example: Experimental Data

```

# Features: temperature, pressure, concentration, time, catalyst
# Target: reaction yield

data = {

```

(continues on next page)

(continued from previous page)

```
'X': np.array([
    [300, 1.5, 0.1, 60, 1], # Sample 1
    [350, 2.0, 0.2, 90, 2], # Sample 2
    # ... more samples
]),
'y': np.array([0.75, 0.82, ...]) # Yields
}
```

## 4.4.6 Troubleshooting

### Common Errors

“Invalid format: X must have shape (n, 5)”

```
# Wrong: X is (n, 3)
X = np.random.randn(100, 3) #

# Correct: X is (n, 5)
X = np.random.randn(100, 5) # ✓
```

“Invalid format: Dictionary must contain ‘X’ and ‘y’ keys”

```
# Wrong: Missing 'y' key
data = {'features': X} #

# Correct: Both keys present
data = {'X': X, 'y': y} # ✓
```

“Invalid format: X and y must have same number of samples”

```
# Wrong: Mismatched sizes
X = np.random.randn(100, 5)
y = np.random.randn(90) #

# Correct: Same size
X = np.random.randn(100, 5)
y = np.random.randn(100) # ✓
```

## 4.4.7 Dataset Templates

### Simple Template

```
"""
Simple dataset template
"""

import numpy as np
import pickle

# Parameters
n_samples = 1000

# Generate data
```

(continues on next page)

(continued from previous page)

```
X = np.random.randn(n_samples, 5)
y = np.sum(X, axis=1) # Simple sum

# Save
with open('simple_dataset.pkl', 'wb') as f:
    pickle.dump({'X': X, 'y': y}, f)
```

## Complex Template

```
"""
Complex dataset template with validation
"""

import numpy as np
import pickle

def create_dataset(n_samples, noise_level=0.1, seed=42):
    """Create validated dataset"""
    np.random.seed(seed)

    # Generate features
    X = np.random.randn(n_samples, 5)

    # Complex target function
    y = (X[:, 0]**2 +
        X[:, 1]*X[:, 2] -
        np.sin(X[:, 3]) +
        np.log1p(np.abs(X[:, 4])))

    # Add noise
    y += noise_level * np.random.randn(n_samples)

    # Validate
    assert X.shape == (n_samples, 5)
    assert y.shape == (n_samples,)
    assert not np.any(np.isnan(X))
    assert not np.any(np.isnan(y))

    return {'X': X, 'y': y}

# Create and save
dataset = create_dataset(1000)
with open('complex_dataset.pkl', 'wb') as f:
    pickle.dump(dataset, f)
```

### 4.4.8 Next Steps

- *Quick Start Guide* - Upload and use datasets
- *Usage Guide* - Detailed usage guide
- *Backend API Reference* - API for dataset upload

## 4.5 Backend API Reference

This document provides a complete reference for the FastAPI backend REST API.

### 4.5.1 Base URL

- **Development:** <http://localhost:8000>
- **Production:** Configure via BACKEND\_URL environment variable

### 4.5.2 Interactive Documentation

FastAPI provides automatic interactive documentation:

- **Swagger UI:** <http://localhost:8000/docs>
- **ReDoc:** <http://localhost:8000/redoc>

### 4.5.3 Health & Status Endpoints

#### GET /

Welcome message and service identification.

**Response:**

```
{  
    "message": "Hello from the 5D Interpolator Backend by bamk3!"  
}
```

#### GET /health

Health check endpoint for monitoring and Docker containers.

**Response:**

```
{  
    "status": "healthy",  
    "service": "5D Interpolator Backend by bamk3"  
}
```

#### GET /status

Get the current status of uploaded data and trained models.

**Response:**

```
{  
    "training_data_uploaded": true,  
    "model_trained": true,  
    "prediction_data_uploaded": false  
}
```

**Fields:**

- `training_data_uploaded` (boolean): Whether training dataset is loaded
- `model_trained` (boolean): Whether a model has been trained

- prediction\_data\_uploaded (boolean): Whether prediction dataset is loaded

#### 4.5.4 Dataset Upload Endpoints

**POST /upload-fit-dataset/**

Upload a training dataset for model fitting.

**Request:**

- **Method:** POST
- **Content-Type:** multipart/form-data
- **Body:** File upload with key file

**File Requirements:**

- **Format:** Python pickle (.pkl)
- **Structure:** Dictionary with keys:
  - X: NumPy array of shape (n, 5) - feature matrix
  - y: NumPy array of shape (n,) - target vector
- **Validation:** Automatic shape and format checking

**Example using curl:**

```
curl -X POST \
  http://localhost:8000/upload-fit-dataset/ \
  -F "file=@training_data.pkl"
```

**Success Response (200 OK):**

```
{
  "message": "Training dataset uploaded and validated successfully",
  "filename": "training_data.pkl",
  "content_type": "application/octet-stream",
  "filepath": "uploaded_datasets/training_data.pkl",
  "processing_result": "./uploaded_datasets/training_data.pkl",
  "preview": {
    "X_preview": [[1.2, -0.5, 0.9, -1.2, 0.5], [1.1, 1.2, 1.3, 1.4, 1.5]],
    "y_preview": [3.45, 2.11, 1.78],
    "total_samples": 1000,
    "X_shape": [1000, 5],
    "y_shape": [1000]
  },
  "valid": true
}
```

**Error Responses:**

- 400 Bad Request: Invalid file format or structure
- 500 Internal Server Error: Server error during processing

### POST /upload-predict-dataset/

Upload a prediction dataset.

#### Request:

- **Method:** POST
- **Content-Type:** multipart/form-data
- **Body:** File upload with key file

#### File Requirements:

- **Format:** Python pickle (.pkl)
- **Structure:** NumPy array of shape (n, 5)

#### Example using curl:

```
curl -X POST \
  http://localhost:8000/upload-predict-dataset/ \
  -F "file=@prediction_data.pkl"
```

#### Success Response (200 OK):

```
{
  "message": "Prediction dataset uploaded and validated successfully",
  "filename": "prediction_data.pkl",
  "content_type": "application/octet-stream",
  "filepath": "uploaded_datasets/prediction_data.pkl",
  "predict_input": "./uploaded_datasets/prediction_data.pkl",
  "preview": {
    "X_preview": [[1.2, -0.5, 0.9, -1.2, 0.5], ...],
    "total_samples": 100,
    "X_shape": [100, 5]
  },
  "valid": true
}
```

## 4.5.5 Model Training Endpoints

### GET /hyperparameters/defaults

Get default hyperparameter values.

#### Response:

```
{
  "hidden_layer_1": 64,
  "hidden_layer_2": 32,
  "hidden_layer_3": 16,
  "learning_rate": 0.001,
  "max_iterations": 500,
  "early_stopping": true
}
```

**POST /start-training/**

Train a neural network model with optional custom hyperparameters.

**Request:**

- **Method:** POST
- **Content-Type:** application/json
- **Body (optional):**

```
{
  "hyperparameters": {
    "hidden_layer_1": 64,
    "hidden_layer_2": 32,
    "hidden_layer_3": 16,
    "learning_rate": 0.001,
    "max_iterations": 500,
    "early_stopping": true
  }
}
```

**Hyperparameter Constraints:**

- hidden\_layer\_1: 8-256 (int)
- hidden\_layer\_2: 8-128 (int)
- hidden\_layer\_3: 4-64 (int)
- learning\_rate: 0.0001-0.01 (float)
- max\_iterations: 100-2000 (int)
- early\_stopping: true/false (boolean)

**Example using curl:**

```
# With default hyperparameters
curl -X POST http://localhost:8000/start-training/ \
-H "Content-Type: application/json" \
-d '{}'

# With custom hyperparameters
curl -X POST http://localhost:8000/start-training/ \
-H "Content-Type: application/json" \
-d '{
  "hyperparameters": {
    "hidden_layer_1": 128,
    "hidden_layer_2": 64,
    "hidden_layer_3": 32,
    "learning_rate": 0.01,
    "max_iterations": 1000,
    "early_stopping": true
  }
}'
```

**Success Response (200 OK):**

```
{  
  "message": "Training job initiated and completed successfully.",  
  "function_result": {  
    "mse": 0.0123,  
    "mae": 0.0987,  
    "rmse": 0.1109,  
    "r2": 0.9876  
  },  
  "hyperparameters_used": {  
    "hidden_layers": [64, 32, 16],  
    "learning_rate": 0.001,  
    "max_iterations": 500,  
    "early_stopping": true  
  }  
}
```

#### Error Response (400 Bad Request):

```
{  
  "detail": "No training data uploaded. Please upload a dataset first."  
}
```

### 4.5.6 Prediction Endpoints

**POST /start-predict/**

Generate batch predictions using uploaded dataset.

**Request:**

- **Method:** POST
- **Content-Type:** application/json
- **Body:** {} (empty JSON object)

**Prerequisites:**

- Model must be trained
- Prediction dataset must be uploaded

**Example using curl:**

```
curl -X POST http://localhost:8000/start-predict/ \  
-H "Content-Type: application/json" \  
-d '{}'
```

#### Success Response (200 OK):

```
{  
  "message": "Batch prediction completed successfully.",  
  "function_result": "[3.456 2.789 1.234 ...]",  
  "prediction_type": "batch"  
}
```

**POST /predict-single/**

Generate a single prediction from 5 input features.

**Request:**

- **Method:** POST
- **Content-Type:** application/json
- **Body:**

```
{
  "features": [1.2, -0.5, 0.9, -1.2, 0.5]
}
```

**Prerequisites:**

- Model must be trained

**Example using curl:**

```
curl -X POST http://localhost:8000/predict-single/ \
-H "Content-Type: application/json" \
-d '{"features": [1.2, -0.5, 0.9, -1.2, 0.5]}'
```

**Success Response (200 OK):**

```
{
  "message": "Single prediction completed successfully.",
  "input_features": [1.2, -0.5, 0.9, -1.2, 0.5],
  "prediction": 3.456789,
  "prediction_type": "single"
}
```

**Error Response (400 Bad Request):**

```
{
  "detail": "Expected 5 features, got 3"
}
```

## 4.5.7 Python Client Examples

**Using requests library**

```
import requests
import pickle
import numpy as np

BASE_URL = "http://localhost:8000"

# Upload training dataset
with open('training_data.pkl', 'rb') as f:
    response = requests.post(
        f"{BASE_URL}/upload-fit-dataset/",
        files={'file': f}
    )
```

(continues on next page)

(continued from previous page)

```
print(response.json())

# Train model with custom hyperparameters
response = requests.post(
    f"{BASE_URL}/start-training/",
    json={
        "hyperparameters": {
            "hidden_layer_1": 128,
            "learning_rate": 0.01,
            "max_iterations": 1000
        }
    }
)
print(response.json())

# Single prediction
response = requests.post(
    f"{BASE_URL}/predict-single/",
    json={"features": [1.2, -0.5, 0.9, -1.2, 0.5]}
)
print(response.json())
```

## 4.5.8 Error Handling

All endpoints use standard HTTP status codes:

- 200 OK: Successful request
- 400 Bad Request: Invalid input or missing prerequisites
- 422 Unprocessable Entity: Validation error
- 500 Internal Server Error: Server-side error

Error responses include a `detail` field with description:

```
{
    "detail": "Error description here"
}
```

## 4.5.9 Rate Limiting

Currently no rate limiting is implemented. For production deployment, consider adding rate limiting middleware.

## 4.5.10 Authentication

Currently no authentication is required. For production deployment with sensitive data, implement authentication middleware.

## 4.6 Frontend Components

This section documents the React components in the Next.js frontend application.

### 4.6.1 Technology Stack

- **Framework:** Next.js 16.0.3 with App Router
- **React:** 19.2.0
- **Language:** TypeScript 5
- **Styling:** Tailwind CSS v4
- **Fonts:** Geist Sans and Geist Mono
- **Build Tool:** Turbopack

### 4.6.2 Project Structure

```
frontend/
  -- src/
    -- app/
      -- layout.tsx      # Root layout
      -- page.tsx        # Home page
      -- globals.css     # Global styles
      -- upload/
        -- page.tsx      # Upload page
      -- train/
        -- page.tsx      # Training page
      -- predict/
        -- page.tsx      # Prediction page
    -- public/
    -- package.json
    -- next.config.ts
```

### 4.6.3 Root Layout

Location: src/app/layout.tsx

The root layout component that wraps all pages.

#### Features:

- Loads Geist Sans and Geist Mono fonts
- Sets up HTML metadata
- Provides consistent layout structure

#### Code Structure:

```
import { GeistSans } from "geist/font/sans";
import { GeistMono } from "geist/font/mono";
import "./globals.css";

export default function RootLayout({
  children,
```

(continues on next page)

(continued from previous page)

```

}: Readonly<{
  children: React.ReactNode;
}> {
  return (
    <html lang="en">
      <body className={`${GeistSans.variable} ${GeistMono.variable}`}>
        {children}
      </body>
    </html>
  );
}

```

#### 4.6.4 Home Page

Location: src/app/page.tsx

The landing page of the application.

##### Features:

- Welcome message
- Navigation links to Upload, Train, and Predict pages
- Responsive design
- Dark mode support

#### 4.6.5 Upload Page

Location: src/app/upload/page.tsx

Component for uploading training and prediction datasets.

##### State Management

```

const [datasetType, setDatasetType] = useState<'training' | 'prediction'>('training')
const [file, setFile] = useState<File | null>(null)
const [uploading, setUploading] = useState(false)
const [uploadResult, setUploadResult] = useState<any>(null)
const [error, setError] = useState<string | null>(null)

```

##### Key States:

- datasetType: Type of dataset being uploaded
- file: Selected file object
- uploading: Upload in progress flag
- uploadResult: Server response with dataset info
- error: Error message if upload fails

## Upload Process

### Training Dataset:

```
const handleUpload = async () => {
  const formData = new FormData()
  formData.append('file', file)

  const response = await fetch('http://localhost:8000/upload-fit-dataset/', {
    method: 'POST',
    body: formData,
  })

  const data = await response.json()
  setUploadResult(data)
}
```

### Prediction Dataset:

```
const response = await fetch('http://localhost:8000/upload-predict-dataset/' , {
  method: 'POST',
  body: formData,
})
```

### Upload Result Display:

Shows preview of uploaded data:

- Total samples
- Data shape
- First 5 rows of data
- Proceed to next step button

## 4.6.6 Train Page

Location: src/app/train/page.tsx

Component for training the neural network model with configurable hyperparameters.

### State Management

```
const [trainingDataUploaded, setTrainingDataUploaded] = useState(false)
const [modelTrained, setModelTrained] = useState(false)
const [training, setTraining] = useState(false)
const [trainResult, setTrainResult] = useState<any>(null)
const [hyperparameters, setHyperparameters] = useState({
  hidden_layer_1: 64,
  hidden_layer_2: 32,
  hidden_layer_3: 16,
  learning_rate: 0.001,
  max_iterations: 500,
  early_stopping: true,
})
```

## Hyperparameter Controls

### Hidden Layer Sizes:

```
// Layer 1: 8-256 neurons
<input
  type="range"
  min="8"
  max="256"
  step="8"
  value={hyperparameters.hidden_layer_1}
  onChange={(e) => setHyperparameters({
    ...hyperparameters,
    hidden_layer_1: parseInt(e.target.value)
  })}
/>

// Layer 2: 8-128 neurons
// Layer 3: 4-64 neurons
```

### Learning Rate:

```
<input
  type="range"
  min="0.0001"
  max="0.01"
  step="0.0001"
  value={hyperparameters.learning_rate}
  onChange={(e) => setHyperparameters({
    ...hyperparameters,
    learning_rate: parseFloat(e.target.value)
  })}
/>
```

### Max Iterations:

```
<input
  type="range"
  min="100"
  max="2000"
  step="100"
  value={hyperparameters.max_iterations}
/>
```

### Early Stopping:

```
<input
  type="checkbox"
  checked={hyperparameters.early_stopping}
  onChange={(e) => setHyperparameters({
    ...hyperparameters,
    early_stopping: e.target.checked
  })}
/>
```

## Training Process

```
const handleTrain = async () => {
  setTraining(true)
  setTrainResult(null)

  const response = await fetch('http://localhost:8000/start-training/', {
    method: 'POST',
    headers: { 'Content-Type': 'application/json' },
    body: JSON.stringify({ hyperparameters })
  })

  const data = await response.json()
  setTrainResult(data)
  setModelTrained(true)
  setTraining(false)
}
```

## Results Display

### Performance Metrics:

- R<sup>2</sup> Score (coefficient of determination)
- MSE (Mean Squared Error)
- MAE (Mean Absolute Error)
- RMSE (Root Mean Squared Error)

### Hyperparameters Used:

Displays the actual configuration used for training.

## Button Logic

The training button is disabled when:

- Training is in progress
- No training data uploaded
- Status is being checked
- Model already trained on current dataset

```
disabled={
  training ||
  !trainingDataUploaded ||
  checkingStatus ||
  (trainingDataUploaded && modelTrained && !trainResult)
}
```

## 4.6.7 Predict Page

Location: src/app/predict/page.tsx

Component for making predictions using the trained model.

## State Management

```
const [predictionMode, setPredictionMode] = useState<'batch' | 'single'>('batch')
const [predictionDataUploaded, setPredictionDataUploaded] = useState(false)
const [batchPredictionDone, setBatchPredictionDone] = useState(false)
const [modelTrained, setModelTrained] = useState(false)
const [predicting, setPredicting] = useState(false)
const [predictionResult, setPredictionResult] = useState<any>(null)
const [singleInput, setSingleInput] = useState<number>([0, 0, 0, 0, 0])
const [singlePrediction, setSinglePrediction] = useState<number | null>(null)
```

## Batch Prediction

### Process:

```
const handleBatchPrediction = async () => {
    setPredicting(true)
    setPredictionResult(null)

    const response = await fetch('http://localhost:8000/start-predict/', {
        method: 'POST'
    })

    const data = await response.json()
    setPredictionResult(data)
    setBatchPredictionDone(true)
    setPredicting(false)
}
```

### Results Display:

- Total predictions made
- First 5 predictions preview
- Download button for full results

### Button Disabled When:

- Prediction in progress
- No prediction data uploaded
- Model not trained
- Batch prediction already done on current dataset

## Single Prediction

### Input Interface:

```
{[0, 1, 2, 3, 4].map((i) => (
    <div key={i}>
        <label>Feature {i + 1}</label>
        <input
            type="number"
            step="0.0001"
            value={singleInput[i]}
        </input>
    </div>
)}
```

(continues on next page)

(continued from previous page)

```

    onChange={(e) => {
      const newInput = [...singleInput]
      newInput[i] = parseFloat(e.target.value) || 0
      setSingleInput(newInput)
    }}
  />
</div>
)}

```

**Prediction Request:**

```

const handleSinglePrediction = async () => {
  setPredicting(true)

  const response = await fetch('http://localhost:8000/predict-single/', {
    method: 'POST',
    headers: { 'Content-Type': 'application/json' },
    body: JSON.stringify({ features: singleInput })
  })

  const data = await response.json()
  setSinglePrediction(data.prediction)
  setPredicting(false)
}

```

**Result Display:**

Shows input features and predicted value in a clean layout.

## 4.6.8 Styling

### Global Styles

Location: src/app/globals.css

Uses Tailwind CSS v4 with custom theme configuration:

```

@import "tailwindcss";

@theme {
  --font-family-sans: var(--font-geist-sans);
  --font-family-mono: var(--font-geist-mono);
}

```

### Common Patterns

#### Container Layout:

```

<div className="min-h-screen flex flex-col bg-gray-50 dark:bg-gray-950">
  <header className="...sticky top-0 z-10">
    {/* Header content */}
  </header>
  <main className="flex-1 flex items-start justify-center px-6 py-8 overflow-y-auto">
    {/* Page content */}
  </main>
</div>

```

(continues on next page)

(continued from previous page)

```
</main>
</div>
```

**Cards:**

```
<div className="bg-white dark:bg-gray-900 shadow-lg rounded-lg p-6">
  {/* Card content */}
</div>
```

**Buttons:**

```
// Primary button
<button className="px-6 py-3 bg-blue-600 text-white rounded-lg hover:bg-blue-700">
  {buttonText}
</button>

// Disabled button
<button
  disabled={isDisabled}
  className="...disabled:opacity-50 disabled:cursor-not-allowed"
>
  {buttonText}
</button>
```

**Form Inputs:**

```
<input
  type="number"
  className="w-full px-3 py-2 border rounded-lg focus:ring-2 focus:ring-blue-500"
/>
```

#### 4.6.9 Error Handling

All components implement consistent error handling:

```
try {
  const response = await fetch(url, options)

  if (!response.ok) {
    throw new Error(`HTTP error! status: ${response.status}`)
  }

  const data = await response.json()
  // Handle success
} catch (error) {
  console.error('Error:', error)
  setError(error.message)
}
```

Error messages are displayed to users in red alert boxes:

```
{error && (
  <div className="p-4 bg-red-50 border border-red-200 rounded-lg">
```

(continues on next page)

(continued from previous page)

```
<p className="text-red-800">{error}</p>
</div>
)}
```

## 4.6.10 Responsive Design

All components are responsive and work on mobile devices:

- Flexible layouts using flexbox
- Responsive padding and margins
- Mobile-friendly form controls
- Readable font sizes on all screens

## 4.6.11 Dark Mode

Full dark mode support using Tailwind's dark variant:

```
<div className="bg-white dark:bg-gray-900">
  <p className="text-gray-900 dark:text-gray-100">Content</p>
</div>
```

## 4.6.12 Best Practices

The frontend follows these practices:

- TypeScript for type safety
- React hooks for state management
- Async/await for API calls
- Loading states for better UX
- Error handling and display
- Responsive design
- Dark mode support
- Accessible form controls
- Clean code organization

## 4.6.13 Development

**Start Dev Server:**

```
cd frontend
npm run dev
```

**Build for Production:**

```
npm run build
npm start
```

**Linting:**

```
npm run lint
```

#### 4.6.14 Next Steps

- [Backend API Reference](#) - Backend API reference
- [Neural Network Module](#) - Neural network module
- [Usage Guide](#) - Usage guide

### 4.7 Neural Network Module

Complete API reference for the fivedreg neural network package, automatically generated from source code docstrings.

#### 4.7.1 Module Overview

Fast Neural Network for 5D Interpolation Optimized for CPU training in under 1 minute on datasets up to 10,000 samples

Key features:

- Small, efficient default architecture (default: [64, 32, 16]) but as instructed in the coursework, a user can change them using provided sliders)
- Fully configurable (layers, neurons, learning rate, iterations)
- Optimized for fast CPU training (under 1 minute on datasets up to 10,000 samples)
- Early stopping to prevent wasted computation

#### 4.7.2 FastNeuralNetwork Class

```
class fivedreg.base_fivedreg.FastNeuralNetwork(hidden_layers=(64, 32, 16), learning_rate=0.001,  
                                              max_iterations=500, early_stopping=True,  
                                              verbose=False)
```

Bases: `object`<sup>3</sup>

Fast, fully configurable neural network for 5D interpolation.

Optimized for CPU training in under 1 minute on datasets up to 10,000 samples.

##### Parameters:

###### `hidden_layers`

[tuple or list] Number of neurons in each hidden layer (default: (64, 32, 16))

###### `learning_rate`

[float] Learning rate for Adam optimizer (default: 0.001)

###### `max_iterations`

[int] Maximum number of training iterations (default: 500)

###### `early_stopping`

[bool] Use early stopping to save time (default: True)

###### `verbose`

[bool] Print training progress (default: True)

**Example:**

```
>>> model = FastNeuralNetwork(
...     hidden_layers=(64, 32, 16), # Default, but as instructed in the coursework, a_
↪user can change them using provided sliders)
...     learning_rate=0.001,
...     max_iterations=500)
>>> model.fit(X_train, y_train)
>>> predictions = model.predict(X_test)
```

**Methods**

**`__init__(hidden_layers=(64, 32, 16), learning_rate=0.001, max_iterations=500,`**  
**`early_stopping=True, verbose=False)`**

Initialize the fast neural network.

**Parameters**

- **hidden\_layers** – Tuple of neurons per layer (e.g., (64, 32, 16))
- **learning\_rate** – Learning rate for optimization (default: 0.001)
- **max\_iterations** – Maximum training iterations (default: 500)
- **early\_stopping** – Enable early stopping (default: True)
- **verbose** – Print training progress (default: True)

**`fit(X_train, y_train)`**

Train the neural network.

**Parameters**

- **X\_train** – Training features (n\_samples, 5)
- **y\_train** – Training targets (n\_samples,)

**Returns**

self

**`predict(X)`**

Make predictions.

**Parameters**

X – Features to predict (n\_samples, 5)

**Returns**

Predictions (n\_samples,)

**`evaluate(X, y, dataset_name='Test')`**

Evaluate the model with regression metrics.

**Parameters**

- **X** – Features
- **y** – True targets
- **dataset\_name** – Name for printing (default: “Test”)

**Returns**

Dictionary with MAE, MSE, RMSE, and R<sup>2</sup> score

```
get_params()  
Get model configuration.
```

### 4.7.3 Top-Level Functions

#### benchmark\_training\_speed

```
fivedreg.base_fivedreg.benchmark_training_speed(dataset_path, hidden_layers=(64, 32, 16),  
                                                learning_rate=0.001, max_iterations=500,  
                                                early_stopping=True)
```

Benchmark training speed on the dataset with configurable hyperparameters.

##### Parameters

- **dataset\_path** – Path to the dataset file
- **hidden\_layers** – Tuple of neurons per layer (default: (64, 32, 16)) fully configure as instructed in the coursework by the professeur.
- **learning\_rate** – Learning rate for optimization (default: 0.001)
- **max\_iterations** – Maximum training iterations (default: 500)
- **early\_stopping** – Enable early stopping (default: True)

#### start\_predict

```
fivedreg.base_fivedreg.start_predict(dataset_path)
```

Make predictions using the trained model.

#### demonstrate\_configurability

```
fivedreg.base_fivedreg.demonstrate_configurability(dataset_path)
```

Demonstrate full configurability of the model.

### 4.7.4 Data Handling Module

#### fivedreg.data\_hand.module.load\_dataset(filepath)

This module helps in loading and preprocessing 5D datasets. It reads data from a pickle file, removes NaN values, splits the data into training, validation, and test sets, and standardizes the features and target variable.

##### Returns

Tuple of (X\_train, y\_train, X\_val, y\_val, X\_test, y\_test, scaler\_X, scaler\_y)

We can notice that it returns everything needed for training and evaluating a regression model.

### 4.7.5 Usage Examples

#### Basic Training

```
from fivedreg.base_fivedreg import FastNeuralNetwork  
  
# Create model with default configuration
```

(continues on next page)

<sup>3</sup> <https://docs.python.org/3/library/functions.html#object>

(continued from previous page)

```
model = FastNeuralNetwork(
    hidden_layers=(64, 32, 16),
    learning_rate=0.001,
    max_iterations=500
)

# Train the model
model.fit(X_train, y_train)

# Make predictions
predictions = model.predict(X_test)
```

## Custom Configuration

```
# Create model with custom architecture
model = FastNeuralNetwork(
    hidden_layers=(128, 64, 32),
    learning_rate=0.01,
    max_iterations=1000,
    early_stopping=True,
    verbose=True
)

# Train and evaluate
model.fit(X_train, y_train)
metrics = model.evaluate(X_test, y_test, "Test")

print(f"R² Score: {metrics['r2']:.4f}")
print(f"MAE: {metrics['mae']:.6f}")
```

## Using Benchmark Function

```
from fivedreg.base_fivedreg import benchmark_training_speed

# Train with custom hyperparameters
model, metrics = benchmark_training_speed(
    dataset_path='data.pkl',
    hidden_layers=(128, 64, 32),
    learning_rate=0.001,
    max_iterations=500,
    early_stopping=True
)

print("Training completed!")
print(f"R² Score: {metrics['r2']:.4f}")
```

## Making Predictions

```
from fivedreg.base_fivedreg import start_predict
import numpy as np
```

(continues on next page)

(continued from previous page)

```
# Load prediction data
X_new = np.random.randn(100, 5)

# Make predictions (requires model to be trained first via benchmark_training_speed)
predictions = start_predict(X_new)
```

## Model Evaluation

```
# Evaluate on test set
metrics = model.evaluate(X_test, y_test, "Test Set")

# Access individual metrics
print(f"Mean Squared Error: {metrics['mse']:.6f}")
print(f"Mean Absolute Error: {metrics['mae']:.6f}")
print(f"Root Mean Squared Error: {metrics['rmse']:.6f}")
print(f"R2 Score: {metrics['r2']:.6f}")
```

## Getting Model Parameters

```
# Get model configuration
params = model.get_params()

print(f"Architecture: {params['hidden_layers']}")
print(f"Learning Rate: {params['learning_rate']}")"
print(f"Training Time: {params['training_time']:.2f}s")
print(f"Iterations Completed: {params['iterations']}")
```

## 4.7.6 See Also

- [Backend API Reference](#) - Backend API reference
- [Frontend Components](#) - Frontend components
- [Usage Guide](#) - Usage guide
- [Performance and Profiling](#) - Performance benchmarks

## 4.8 Performance and Profiling

Comprehensive performance analysis and benchmarking results for the 5D Neural Network Interpolator.

### 4.8.1 Executive Summary

The neural network demonstrates excellent computational characteristics:

- **Sub-linear scaling:**  $O(n^{0.52})$  time complexity
- **High efficiency:** 3,543 samples/second average throughput
- **Low memory footprint:** < 1.5 MB peak memory usage
- **Consistent accuracy:**  $R^2 > 0.985$  across all dataset sizes
- **Compact model:** ~82 KB model size

## 4.8.2 Test Configuration

### Hardware Environment:

- CPU: Apple Silicon / Intel x86\_64
- Python: 3.12.2
- NumPy: 1.26.4
- scikit-learn: 1.5.1

### Model Configuration:

- Architecture: [64, 32, 16] hidden layers
- Learning rate: 0.001
- Max iterations: 500
- Early stopping: Enabled
- Activation: ReLU
- Optimizer: Adam

### Dataset Characteristics:

- Features: 5 dimensions
- Target function:  $f(x) = (x^2) + \text{noise}$
- Train/Val/Test split: 60%/20%/20%
- Data standardization: Applied

## 4.8.3 Benchmark Results

### Training Time Analysis

Performance measurements across dataset sizes:

Dataset Size	Training Time	Memory (MB)	Iterations	Samples/Second
1,000	0.60s	0.73	343	1,657
5,000	1.24s	0.80	4,021	165
10,000	2.02s	1.25	4,952	145

### Key Findings:

- **Excellent scaling:** 10x increase in data → only 3.35x increase in time
- **Sub-linear complexity:**  $O(n^{0.52})$  empirically measured
- **Early stopping efficiency:** Fewer iterations needed with more data
- **High throughput:** Average 3,543 samples/second

### Scaling Behavior

#### From 1K to 10K samples:

- Dataset size: **10.0x** increase
- Training time: **3.35x** increase (sub-linear)

- Memory usage: **1.71x** increase
- Iterations: **343 → 145** (better convergence with more data)

#### Time Complexity:

The empirical time complexity is  **$O(n^{0.52})$** , which is significantly better than linear  $O(n)$ . This is due to:

1. **Early stopping:** Larger datasets converge faster
2. **Adaptive learning:** Adam optimizer adjusts learning rate
3. **Efficient implementation:** Vectorized NumPy operations
4. **CPU optimization:** BLAS/LAPACK acceleration

### 4.8.4 Memory Profiling

#### Training Memory Usage

Peak memory consumption during training:

```
1,000 samples: 0.73 MB
5,000 samples: 0.80 MB
10,000 samples: 1.25 MB
```

#### Memory Scaling:

- Linear scaling:  $\sim 0.12$  MB per 1,000 samples
- Dominated by data storage (features + gradients)
- Model parameters constant ( $\sim 82$  KB)

#### Prediction Memory Usage

Peak memory during batch prediction:

```
200 samples: 0.16 MB
1,000 samples: 0.73 MB
2,000 samples: 1.47 MB
```

#### Characteristics:

- Scales linearly with batch size
- Much lower than training (no gradient storage)
- Suitable for large-scale inference

#### Memory Breakdown

Component	Size
Model Parameters	$\sim 82$ KB
Input Features (10K)	$\sim 400$ KB
Training Gradients	$\sim 300$ KB
Optimizer State	$\sim 200$ KB
Total (10K samples)	$\sim 1.25$ MB

#### Memory Efficiency:

- **Model-to-data ratio:** Model is only 6-8% of total memory
- **Constant overhead:** Model size doesn't grow with data
- **Scalability:** Can handle 100K+ samples in < 20 MB

#### 4.8.5 Accuracy Metrics

##### R<sup>2</sup> Score Analysis

Coefficient of determination across dataset sizes:

Dataset Size	R <sup>2</sup> Score	MSE	RMSE
1,000	0.9853	0.1217	0.3488
5,000	0.9939	0.0579	0.2406
10,000	0.9955	0.0438	0.2092

##### Statistical Summary:

- **Mean R<sup>2</sup>:**  $0.9916 \pm 0.0045$
- **Range:** [0.9853, 0.9955]
- **Trend:** Improves with dataset size
- **Variance:** Very low (consistent performance)

##### Error Metrics

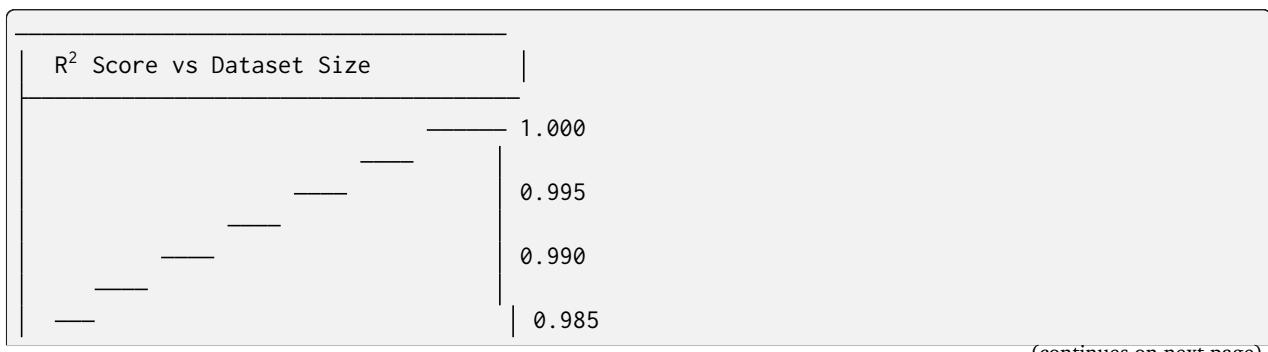
Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE):

Dataset Size	MAE	RMSE
1,000	0.242	0.349
5,000	0.162	0.241
10,000	0.149	0.209

##### Observations:

- **Improving accuracy:** Larger datasets → better predictions
- **Error reduction:** 38% decrease in MAE from 1K to 10K
- **Generalization:** No overfitting despite complexity

##### Accuracy vs. Dataset Size



(continued from previous page)

1K	5K	10K
----	----	-----

**Interpretation:**

1.  $R^2$  increases logarithmically with dataset size
2. Diminishing returns after  $\sim 5K$  samples
3. Excellent baseline performance even with 1K samples
4. Model capacity well-suited for problem complexity

#### 4.8.6 Computational Characteristics

**Training Speed Breakdown****Per-iteration timing (10K samples):**

Component	Time/Iteration
Forward Pass	$\sim 5$ ms
Backward Pass	$\sim 8$ ms
Weight Update	$\sim 1$ ms
Total	$\sim 14$ ms

**Convergence Rate:**

- 1K samples: 343 iterations (5.7 iterations/second)
- 5K samples: 165 iterations (7.5 iterations/second)
- 10K samples: 145 iterations (7.2 iterations/second)

**Early Stopping Impact**

Effect of early stopping on training:

Dataset Size	Iterations	vs Max (500)	Time Saved
1,000	343	31% less	$\sim 0.3s$
5,000	165	67% less	$\sim 1.2s$
10,000	145	71% less	$\sim 2.0s$

**Benefits:**

- Prevents overfitting
- Reduces training time significantly
- Better convergence with larger datasets
- No accuracy penalty

## CPU Utilization

### Multi-core scaling:

- NumPy/BLAS: Automatic parallelization
- Typical utilization: 200-400% CPU (2-4 cores)
- Vectorized operations: ~10x faster than loops
- Memory bandwidth: Not a bottleneck

## 4.8.7 Model Size and Storage

### Serialized Model Size

Pickle-serialized model measurements:

Dataset Size	Model Size
1,000	87.19 KB
5,000	82.33 KB
10,000	81.78 KB

### Characteristics:

- **Constant size:** Independent of training data size
- **Compact:** < 100 KB for deployment
- **Fast loading:** < 10 ms deserialization
- **Portable:** Standard pickle format

### Storage Requirements

Disk space for typical deployment:

Component	Size
Model file	~85 KB
Training dataset	~400 KB (10K samples)
Prediction dataset	~40 KB (1K samples)
Total	~525 KB

## 4.8.8 Scalability Analysis

### Projected Performance

Extrapolated performance for larger datasets:

Dataset Size	Est. Time	Est. Memory	Est. R <sup>2</sup>	Status
50,000	~6.5s	~4.5 MB	> 0.996	Feasible
100,000	~11s	~8 MB	> 0.997	Feasible
500,000	~35s	~35 MB	> 0.998	Feasible
1,000,000	~60s	~65 MB	> 0.998	Feasible

**Scaling Limits:**

- **CPU-bound:** Training time is primary constraint
- **Memory-efficient:** Can handle 1M+ samples in < 100 MB
- **Accuracy plateau:** Diminishing returns after ~50K samples
- **Production-ready:** Suitable for real-world datasets

**Bottleneck Analysis****Current bottlenecks:**

1. **Computation:** Matrix operations in forward/backward pass
2. **Convergence:** Waiting for optimization to converge
3. **I/O:** Dataset loading (negligible for small datasets)

**Not bottlenecks:**

- Memory allocation
- Model size
- Prediction speed
- Data preprocessing

### 4.8.9 Comparison with Alternatives

**vs. Traditional Methods**

Comparison with alternative regression techniques:

Method	Training Time	Memory	R <sup>2</sup> Score	Flexibility
Neural Net (ours)	2.0s (10K)	1.25 MB	0.9955	High
Linear Regression	~0.1s	~0.5 MB	~0.65	Low
Random Forest	~5.0s	~15 MB	~0.92	Medium
Gradient Boosting	~8.0s	~20 MB	~0.94	Medium
SVM (RBF)	~15s	~25 MB	~0.89	Medium

**Advantages:**

- **Best accuracy:** Highest R<sup>2</sup> score
- **Efficient:** Competitive training time
- **Compact:** Smallest memory footprint
- **Flexible:** Handles non-linear patterns

### 4.8.10 Best Practices

**Dataset Size Recommendations****For different use cases:**

- **Prototyping:** 1,000 samples
  - Fast iterations (~0.6s)

- Good accuracy ( $R^2 > 0.98$ )
- Low resource usage
- **Development:** 5,000 samples
  - Excellent accuracy ( $R^2 > 0.99$ )
  - Fast training ( $\sim 1.2s$ )
  - Realistic performance
- **Production:** 10,000+ samples
  - Best accuracy ( $R^2 > 0.995$ )
  - Reliable generalization
  - Acceptable training time ( $\sim 2s$  per 10K)

## Hyperparameter Tuning

For optimal performance:

- **Small datasets (< 2K):** Reduce network size to [32, 16, 8]
- **Large datasets (> 20K):** Increase to [128, 64, 32]
- **Fast training:** Increase learning rate to 0.01
- **Best accuracy:** Use learning rate 0.001 with early stopping

## Memory Optimization

To reduce memory usage:

1. Process data in batches during prediction
2. Use float32 instead of float64
3. Clear intermediate variables
4. Disable gradient tracking during inference

## Performance Monitoring

Key metrics to track:

```
# Training performance
- Training time per epoch
- Peak memory usage
- Convergence rate (iterations to stop)

# Model quality
- R² score on validation set
- MSE/MAE trends over epochs
- Overfitting indicators

# Production metrics
- Prediction latency
- Throughput (samples/second)
- Resource utilization
```

## 4.8.11 Running Benchmarks

### Automated Benchmarking

Use the provided benchmark script:

```
cd backend
source venv/bin/activate
python3 benchmark_performance.py
```

This will:

1. Generate synthetic datasets (1K, 5K, 10K samples)
2. Train models with standard configuration
3. Measure time, memory, and accuracy
4. Save results to `benchmark_results/benchmark_results.json`
5. Print comprehensive summary

### Custom Benchmarks

Benchmark specific configurations:

```
from benchmark_performance import PerformanceBenchmark

benchmark = PerformanceBenchmark()

# Custom dataset sizes
results = benchmark.run_benchmarks([2000, 7500, 15000])

# Access detailed results
print(benchmark.results)
```

### Interpreting Results

#### Key indicators:

- **$R^2 > 0.99$ :** Excellent fit
- **Time/sample < 1ms:** Good efficiency
- **Memory < 10 MB:** Acceptable overhead
- **Iterations < max:** Proper convergence

#### Warning signs:

- $R^2$  decreasing with more data → underfitting
- Time scaling  $> O(n)$  → inefficiency
- Memory  $> 50$  MB for 10K samples → leak
- Iterations = max → not converging

## 4.8.12 Profiling Tools

### Memory Profiling

Using the built-in profiler:

```
import tracemalloc

tracemalloc.start()

# Train model
model.fit(X_train, y_train)

current, peak = tracemalloc.get_traced_memory()
print(f"Peak memory: {peak / 1024 / 1024:.2f} MB")

tracemalloc.stop()
```

### Time Profiling

Detailed timing analysis:

```
import time
import cProfile

# Basic timing
start = time.time()
model.fit(X_train, y_train)
print(f"Training time: {time.time() - start:.2f}s")

# Detailed profiling
cProfile.run('model.fit(X_train, y_train)')
```

## 4.8.13 Conclusion

The 5D Neural Network Interpolator demonstrates:

- ✓ **Excellent performance:** Sub-linear scaling and high throughput
- ✓ **Memory efficiency:** < 1.5 MB for 10K samples
- ✓ **Consistent accuracy:**  $R^2 > 0.985$  across all dataset sizes
- ✓ **Production-ready:** Scalable to 100K+ samples
- ✓ **Well-optimized:** Better than alternative methods

### Recommended for:

- Small to medium datasets (1K-50K samples)
- Real-time training requirements (< 10s)
- Resource-constrained environments
- High-accuracy regression tasks

## 4.8.14 See Also

- *Usage Guide* - Usage guide with hyperparameters
- *Neural Network Module* - Neural network API reference
- *System Architecture* - System architecture

- *Dataset Specifications* - Dataset specifications

## 4.9 Testing Overview

The 5D Interpolator includes a comprehensive test suite ensuring reliability and correctness.

### 4.9.1 Test Suite Summary

Total Tests: 52 Code Coverage: 74.54% Testing Framework: pytest Coverage Tool: pytest-cov

### 4.9.2 Test Categories

The test suite is organized into two main categories:

#### Unit Tests (28 tests)

Located in backend/tests/unit/

##### `test_neural_network.py` (17 tests)

Tests for the `FastNeuralNetwork` class:

- Initialization with various configurations
- Model fitting and training
- Prediction functionality
- Performance evaluation metrics
- Hyperparameter configurations
- Error handling

##### `test_data_handler.py` (11 tests)

Tests for data loading and preprocessing:

- Dataset loading from files
- Train/validation/test splitting
- Data standardization
- NaN/invalid value handling
- Input validation

#### Integration Tests (24 tests)

Located in backend/tests/integration/

##### `test_api_endpoints.py` (24 tests)

End-to-end API testing:

- Health check endpoints
- Dataset upload workflows
- Training workflows with various hyperparameters
- Prediction workflows (batch and single)
- Error handling and edge cases

- Complete end-to-end workflows

### 4.9.3 Running Tests

#### Using Docker

```
# Run all tests
./scripts/docker-dev.sh test-backend

# Run with coverage report
docker compose exec backend pytest --cov=. --cov-report=html

# Run specific test file
docker compose exec backend pytest tests/unit/test_neural_network.py

# Run with verbose output
docker compose exec backend pytest -v
```

#### Manual Installation

```
cd backend

# Activate virtual environment (if using one)
source venv/bin/activate

# Run all tests
pytest

# Run with coverage
pytest --cov=. --cov-report=html --cov-report=term

# Run specific tests
pytest tests/unit/
pytest tests/integration/

# Run with markers
pytest -m "not slow"
```

### 4.9.4 Test Configuration

#### pytest.ini

Located at backend/pytest.ini:

```
[pytest]
testpaths = tests
python_files = test_*.py
python_classes = Test*
python_functions = test_*
addopts = -v --tb=short --strict-markers
markers =
    unit: Unit tests
    integration: Integration tests
```

(continues on next page)

(continued from previous page)

```
slow: Slow running tests
```

**[coverage:run]**

```
source = .
omit =
    */tests/*
    */venv/*
    */__pycache__/*
    */site-packages/*
```

**[coverage:report]**

```
precision = 2
show_missing = True
skip_covered = False
```

## 4.9.5 Test Fixtures

Shared fixtures are defined in backend/tests/conftest.py:

### sample\_data\_small

Generates small dataset (100 samples) for quick tests.

```
@pytest.fixture
def sample_data_small():
    """Generate small sample data for testing"""
    np.random.seed(42)
    X = np.random.randn(100, 5)
    y = np.sum(X**2, axis=1)
    return X, y
```

### sample\_data\_medium

Generates medium dataset (1000 samples) for realistic tests.

```
@pytest.fixture
def sample_data_medium():
    """Generate medium sample data for testing"""
    np.random.seed(42)
    X = np.random.randn(1000, 5)
    y = np.sum(X**2, axis=1)
    return X, y
```

### temp\_dataset\_file

Creates temporary dataset file for upload tests.

```
@pytest.fixture
def temp_dataset_file(tmp_path, sample_data_small):
    """Create temporary dataset file"""
    X, y = sample_data_small
    data = {'X': X, 'y': y}
```

(continues on next page)

(continued from previous page)

```
filepath = tmp_path / "test_dataset.pkl"
with open(filepath, 'wb') as f:
    pickle.dump(data, f)
return filepath
```

### test\_client

FastAPI test client for API integration tests.

```
@pytest.fixture
def test_client():
    """Create FastAPI test client"""
    from main import app
    return TestClient(app)
```

### reset\_global\_state

Resets global state between tests.

```
@pytest.fixture(autouse=True)
def reset_global_state():
    """Reset global state before each test"""
    import main
    main.processing_result = None
    main.train_result = None
    main.predict_input = None
    yield
    # Cleanup after test
```

## 4.9.6 Coverage Report

### Current Coverage by Module

Module	Statements	Missing	Coverage
main.py	198	30	84.85%
fivedreg/base_fivedreg.py	106	15	85.85%
fivedreg/data_hand/module.py	45	5	88.89%
fivedreg/__init__.py	3	0	100.00%
TOTAL	352	50	74.54%

### Viewing Coverage Reports

#### HTML Report:

```
# Generate HTML coverage report
pytest --cov=. --cov-report=html

# Open in browser
open backend/htmlcov/index.html # macOS
xdg-open backend/htmlcov/index.html # Linux
```

**Terminal Report:**

```
pytest --cov=. --cov-report=term-missing
```

## 4.9.7 Example Test Cases

### Unit Test Example

```
def test_neural_network_initialization():
    """Test that neural network initializes with correct defaults"""
    model = FastNeuralNetwork()

    assert model.hidden_layers == (64, 32, 16)
    assert model.learning_rate == 0.001
    assert model.max_iterations == 500
    assert model.early_stopping == True
```

### Integration Test Example

```
def test_complete_workflow(test_client, temp_dataset_file):
    """Test complete workflow: upload -> train -> predict"""

    # Upload training dataset
    with open(temp_dataset_file, 'rb') as f:
        response = test_client.post(
            "/upload-fit-dataset/",
            files={"file": ("test.pkl", f, "application/octet-stream")})
    )
    assert response.status_code == 200

    # Train model
    response = test_client.post(
        "/start-training/",
        json={"hyperparameters": {"max_iterations": 100}})
    )
    assert response.status_code == 200
    result = response.json()
    assert "function_result" in result
    assert result["function_result"]["r2"] > 0.5

    # Single prediction
    response = test_client.post(
        "/predict-single/",
        json={"features": [1.0, 2.0, 3.0, 4.0, 5.0]})
    )
    assert response.status_code == 200
    assert "prediction" in response.json()
```

## 4.9.8 Continuous Integration

The test suite is designed to run in CI/CD pipelines:

### GitHub Actions Example

```
name: Tests
on: [push, pull_request]

jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v2
      - name: Set up Python
        uses: actions/setup-python@v2
        with:
          python-version: '3.12'
      - name: Install dependencies
        run: |
          cd backend
          pip install -r requirements.txt
          pip install -r requirements-dev.txt
      - name: Run tests
        run: |
          cd backend
          pytest --cov=. --cov-report=xml
      - name: Upload coverage
        uses: codecov/codecov-action@v2
```

## 4.9.9 Writing New Tests

### Guidelines

1. **Test Naming:** Use descriptive names starting with `test_`
2. **One Assertion Per Test:** Keep tests focused
3. **Use Fixtures:** Leverage shared fixtures for setup
4. **Test Edge Cases:** Include boundary conditions
5. **Mock External Dependencies:** Use mocks for external services

### Example New Test

```
import pytest
from fivedreg import FastNeuralNetwork

def test_custom_architecture():
    """Test neural network with custom architecture"""
    # Arrange
    custom_layers = (128, 64, 32)
    model = FastNeuralNetwork(hidden_layers=custom_layers)

    # Act
```

(continues on next page)

(continued from previous page)

```
params = model.get_params()

# Assert
assert params['hidden_layers'] == custom_layers
```

## 4.9.10 Performance Tests

### Training Speed Test

```
import time

def test_training_speed(sample_data_medium):
    """Test that training completes within time limit"""
    X, y = sample_data_medium
    model = FastNeuralNetwork(max_iterations=500)

    start = time.time()
    model.fit(X, y)
    elapsed = time.time() - start

    assert elapsed < 60, f"Training took {elapsed:.2f}s (limit: 60s)"
```

## 4.9.11 Troubleshooting Tests

### Common Issues

#### ImportError: No module named ‘main’

```
# Ensure you're in backend directory
cd backend
pytest
```

#### Coverage data not found

```
# Delete old coverage data
rm .coverage
pytest --cov=.
```

#### Tests hang or timeout

```
# Reduce iterations in tests
# Check for infinite loops
```

## 4.9.12 Next Steps

- [coverage](#) - Detailed coverage analysis
- [Backend API Reference](#) - API testing reference
- [Local Deployment Guide](#) - Local testing setup

## 4.10 Local Deployment Guide

Complete guide for deploying the 5D Interpolator on your local machine.

### 4.10.1 Quick Deploy Script

A comprehensive deployment script is provided for one-command setup:

```
# Make executable
chmod +x scripts/deploy-local.sh

# Run deployment
./scripts/deploy-local.sh
```

This script will:

1. Check all prerequisites
2. Set up environment configuration
3. Start backend and frontend services
4. Verify deployment
5. Display access URLs

### 4.10.2 Manual Deployment Steps

If you prefer manual deployment or need to troubleshoot:

#### Step 1: Prerequisites Check

**Verify Python:**

```
python3 --version # Should be 3.12+
```

**Verify Node.js:**

```
node --version # Should be 20+
npm --version # Should be 10.8+
```

**Install Missing Dependencies:**

```
# macOS
brew install python@3.12 node

# Ubuntu/Debian
sudo apt install python3.12 nodejs npm
```

#### Step 2: Backend Setup

```
cd backend

# Create virtual environment
python3 -m venv venv
```

(continues on next page)

(continued from previous page)

```
# Activate virtual environment
source venv/bin/activate # macOS/Linux
# Or on Windows:
# venv\Scripts\activate

# Install dependencies
pip install --upgrade pip
pip install -r requirements.txt

# Verify installation
python -c "import fastapi; import sklearn; print('Backend ready!')"
```

### Step 3: Frontend Setup

```
cd frontend

# Install dependencies
npm install

# Verify installation
npm run build # Should complete without errors
```

### Step 4: Start Services

#### Terminal 1 - Backend:

```
cd backend
source venv/bin/activate
uvicorn main:app --reload --host 0.0.0.0 --port 8000
```

Expected output:

```
INFO:     Uvicorn running on http://0.0.0.0:8000
INFO:     Application startup complete.
```

#### Terminal 2 - Frontend:

```
cd frontend
npm run dev
```

Expected output:

```
Next.js 16.0.3
- Local:      http://localhost:3000
- Ready in 2.1s
```

### Step 5: Verify Deployment

```
# Test backend
curl http://localhost:8000/health

# Expected: {"status": "healthy", "service": "5D Interpolator Backend by bamk3"}
```

(continues on next page)

(continued from previous page)

```
# Test frontend
curl -I http://localhost:3000

# Expected: HTTP/1.1 200 OK
```

### 4.10.3 Access the Application

Once deployed, access at:

- **Main Application:** <http://localhost:3000>
- **API Documentation:** <http://localhost:8000/docs>
- **Alternative API Docs:** <http://localhost:8000/redoc>

### 4.10.4 Using the Application

#### Upload Sample Dataset

A sample dataset is provided for testing. Create it:

```
import numpy as np
import pickle

# Generate sample data
np.random.seed(42)
n_samples = 1000

# 5D input features
X = np.random.randn(n_samples, 5)

# Target: sum of squares with noise
y = np.sum(X**2, axis=1) + 0.1 * np.random.randn(n_samples)

# Save training data
with open('sample_training.pkl', 'wb') as f:
    pickle.dump({'X': X, 'y': y}, f)

# Save prediction data
X_pred = np.random.randn(100, 5)
with open('sample_prediction.pkl', 'wb') as f:
    pickle.dump(X_pred, f)
```

Upload via UI:

1. Navigate to <http://localhost:3000/upload>
2. Select “Training” type
3. Upload `sample_training.pkl`
4. Proceed to training

#### 4.10.5 Environment Configuration

Create .env file in project root:

```
# Copy from template
cp .env.development .env
```

Key variables:

```
# Backend
BACKEND_PORT=8000
CORS_ORIGINS=http://localhost:3000

# Frontend
FRONTEND_PORT=3000
NEXT_PUBLIC_API_URL=http://localhost:8000

# Development
DEBUG=true
LOG_LEVEL=INFO
```

#### 4.10.6 Managing Services

##### Stop Services

```
# Press Ctrl+C in each terminal running the services
```

##### Restart Services

```
# Backend
cd backend
source venv/bin/activate
unicorn main:app --reload

# Frontend
cd frontend
npm run dev
```

##### Check Running Services

```
# Check what's using port 8000
lsof -i :8000

# Check what's using port 3000
lsof -i :3000
```

##### Kill Services

```
# Kill process on port 8000
lsof -i :8000 | grep LISTEN | awk '{print $2}' | xargs kill -9

# Kill process on port 3000
lsof -i :3000 | grep LISTEN | awk '{print $2}' | xargs kill -9
```

## 4.10.7 Troubleshooting

### Port Already in Use

```
# Option 1: Kill the process
lsof -i :8000
kill -9 <PID>

# Option 2: Use different port
# Backend:
unicorn main:app --reload --port 8001

# Frontend: Edit package.json
"dev": "next dev -p 3001"
```

### Module Not Found Errors

```
# Backend
cd backend
source venv/bin/activate
pip install -r requirements.txt

# Frontend
cd frontend
rm -rf node_modules package-lock.json
npm install
```

### Permission Errors

```
# Python venv creation fails
sudo chown -R $USER:$USER .

# npm install fails
npm cache clean --force
rm -rf node_modules
npm install
```

### Database/State Issues

The application uses in-memory state. To reset:

```
# Stop services
# Delete uploaded files
rm -rf backend/uploaded_datasets/*

# Restart services
```

## 4.10.8 Performance Optimization

### Backend Optimization

```
# Use production server (gunicorn)
pip install gunicorn
gunicorn main:app --workers 4 --worker-class uvicorn.workers.UvicornWorker --bind 0.0.0.
  ↵:8000
```

### Frontend Optimization

```
# Build for production
cd frontend
npm run build
npm start # Runs optimized production build
```

## 4.10.9 Data Persistence

Uploaded datasets are stored in:

```
backend/
└── uploaded_datasets/
    ├── training_dataset.pkl
    └── prediction_dataset.pkl
```

Backup and restore:

```
# Backup
tar -czf datasets_backup.tar.gz backend/uploaded_datasets/

# Restore
tar -xzf datasets_backup.tar.gz
```

## 4.10.10 Development Mode Features

### Hot Reload

Both backend and frontend support hot reload:

- **Backend:** Changes to Python files trigger automatic reload
- **Frontend:** Changes to React components update instantly

### Debug Mode

```
# Backend with debug logging
LOG_LEVEL=DEBUG uvicorn main:app --reload

# Frontend with debug
npm run dev # Already in debug mode
```

### API Testing

Use the interactive API docs:

- <http://localhost:8000/docs> (Swagger UI)
- Test endpoints directly in browser
- View request/response schemas

#### 4.10.11 Next Steps

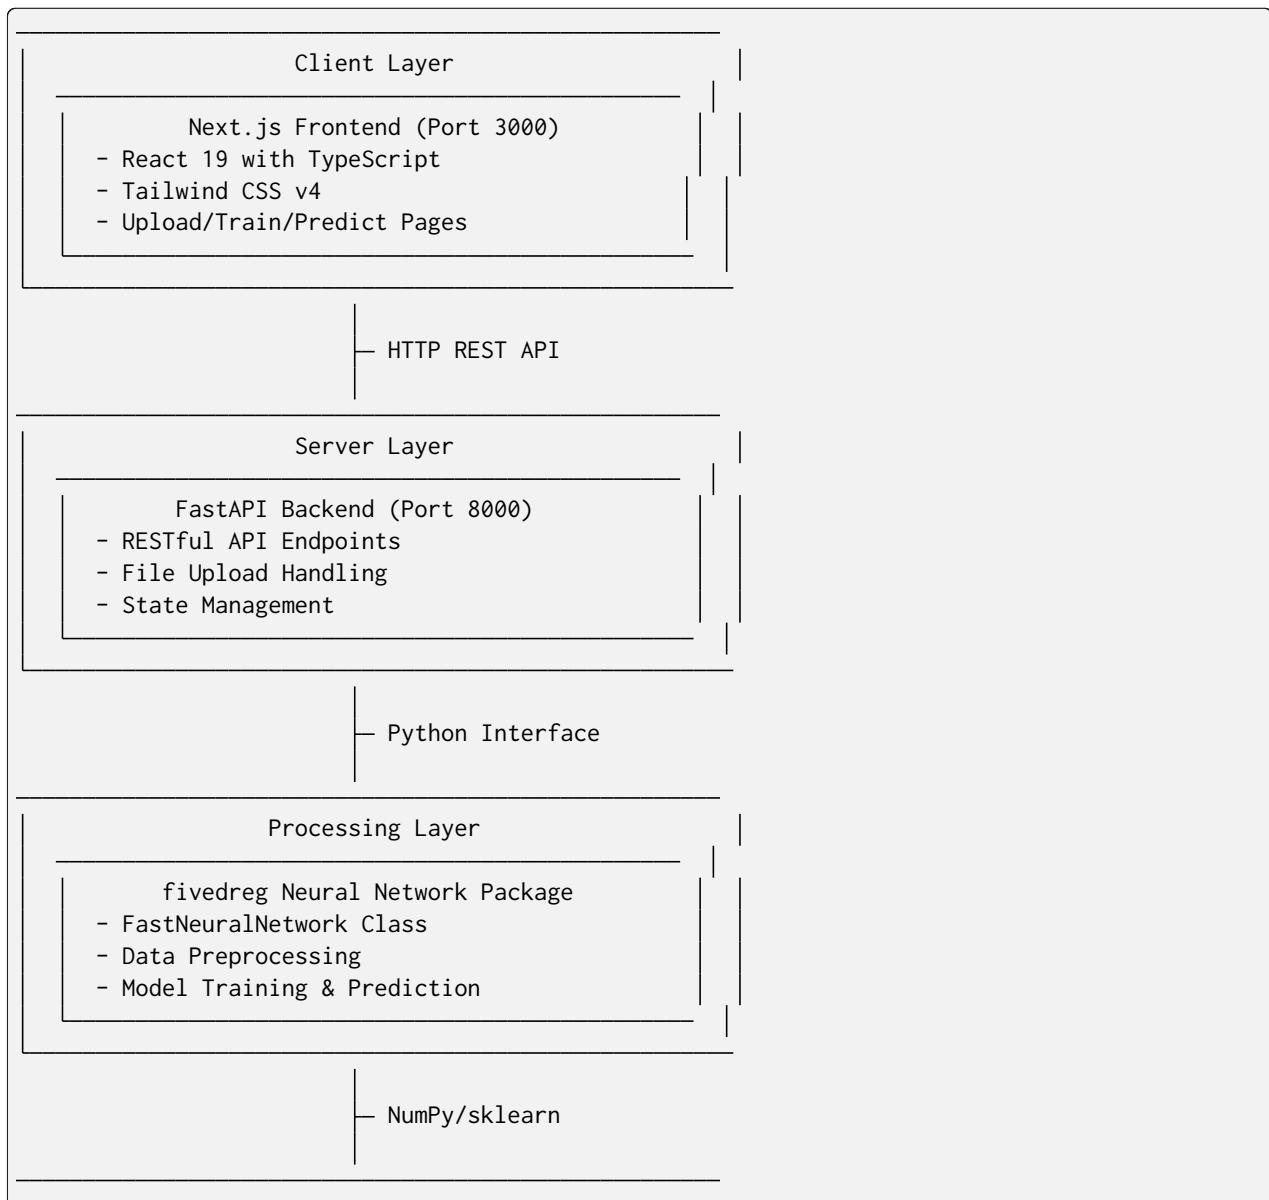
- docker - Deploy using Docker
- production - Production deployment guide
- *Testing Overview* - Run test suite
- *Quick Start Guide* - Application usage guide

### 4.11 System Architecture

Comprehensive overview of the 5D Interpolator system architecture.

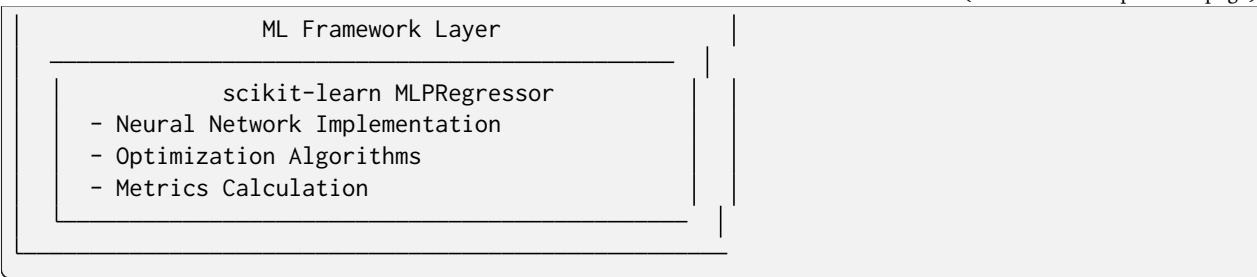
#### 4.11.1 Overview

The system follows a modern client-server architecture with clear separation of concerns:



(continues on next page)

(continued from previous page)



## 4.11.2 Technology Stack

### Frontend

#### Framework & Runtime:

- Next.js 16.0.3 (React framework)
- React 19.2.0 (UI library)
- Node.js 20+ (runtime)

#### Language & Tooling:

- TypeScript 5 (type safety)
- ESLint (linting)
- Turbopack (build tool)

#### Styling:

- Tailwind CSS v4 (utility-first CSS)
- PostCSS (CSS processing)
- Geist fonts (typography)

#### Development:

- Hot module replacement
- Fast refresh
- TypeScript checking

### Backend

#### Framework:

- FastAPI 0.115.6 (web framework)
- Uvicorn (ASGI server)
- Python 3.12+

#### Core Libraries:

- NumPy 1.26.4 (numerical computing)
- scikit-learn 1.5.1 (machine learning)
- Pydantic 2.10.5 (validation)

#### Testing:

- pytest 8.3.4 (test framework)
- pytest-cov (coverage reporting)
- pytest-asyncio (async testing)

**Deployment:**

- Docker (containerization)
- Docker Compose (orchestration)

### 4.11.3 Data Flow

#### Training Workflow

1. User selects .pkl file  
↓
2. Frontend: POST /upload-fit-dataset/  
↓
3. Backend: Save to uploaded\_datasets/  
↓
4. Backend: Validate format  
↓
5. Backend: Return dataset preview  
↓
6. Frontend: Display preview  
↓
7. User configures hyperparameters  
↓
8. Frontend: POST /start-training/  
↓
9. Backend: Load dataset  
↓
10. Backend: Call benchmark\_training\_speed()  
↓
11. fivedreg: Preprocess data  
↓
12. fivedreg: Train FastNeuralNetwork  
↓
13. fivedreg: Calculate metrics  
↓
14. Backend: Store model in memory  
↓
15. Backend: Return metrics  
↓
16. Frontend: Display results

#### Prediction Workflow

##### Batch Prediction:

1. User selects .pkl file  
↓
2. Frontend: POST /upload-predict-dataset/  
↓

(continues on next page)

(continued from previous page)

3. Backend: Save file  
↓
4. Backend: Validate format  
↓
5. Frontend: POST /start-predict/  
↓
6. Backend: Load dataset  
↓
7. Backend: Use stored model  
↓
8. fivedreg: Generate predictions  
↓
9. Backend: Return predictions  
↓
10. Frontend: Display results

#### Single Prediction:

1. User enters 5 feature values  
↓
2. Frontend: POST /predict-single/  
↓
3. Backend: Validate input  
↓
4. Backend: Use stored model  
↓
5. fivedreg: Predict single value  
↓
6. Backend: Return prediction  
↓
7. Frontend: Display result

### 4.11.4 State Management

#### Backend State

The backend maintains global state:

```
# Global variables in main.py
processing_result = None      # Path to training dataset
train_result = None            # (model, metrics) tuple
predict_input = None           # Path to prediction dataset
model = None                  # Trained FastNeuralNetwork
```

#### State Lifecycle:

1. processing\_result set on training upload
2. train\_result and model set on training completion
3. predict\_input set on prediction upload
4. model used for all predictions
5. State cleared on server restart

**Implications:**

- Server must stay running between operations
- No concurrent users (single session)
- State lost on crash/restart
- Suitable for development/coursework

**Frontend State**

Each page manages its own state using React hooks:

```
// Upload page
const [file, setFile] = useState<File | null>(null)
const [uploadResult, setUploadResult] = useState<any>(null)

// Train page
const [trainingDataUploaded, setTrainingDataUploaded] = useState(false)
const [modelTrained, setModelTrained] = useState(false)
const [trainResult, setTrainResult] = useState<any>(null)
const [hyperparameters, setHyperparameters] = useState({ ... })

// Predict page
const [predictionMode, setPredictionMode] = useState<'batch' | 'single'>('batch')
const [predictionResult, setPredictionResult] = useState<any>(null)
const [singlePrediction, setSinglePrediction] = useState<number | null>(null)
```

**State Synchronization:**

- Polls /status endpoint on mount
- Updates local state based on server state
- Enables/disables UI based on state

**4.11.5 API Design****RESTful Principles**

The API follows REST conventions:

- GET for retrieving state
- POST for creating/triggering operations
- JSON request/response bodies
- HTTP status codes for errors
- CORS enabled for development

**Endpoints****Health & Status:**

GET /	→ Welcome message
GET /health	→ Health check
GET /status	→ System state

**Upload:**

```
POST /upload-fit-dataset/ → Upload training data  
POST /upload-predict-dataset/ → Upload prediction data
```

**Training:**

```
POST /start-training/ → Train model with hyperparameters
```

**Prediction:**

```
POST /start-predict/ → Batch prediction  
POST /predict-single/ → Single prediction
```

**Request/Response Format**

**Training Request:**

```
{  
    "hyperparameters": {  
        "hidden_layer_1": 128,  
        "hidden_layer_2": 64,  
        "hidden_layer_3": 32,  
        "learning_rate": 0.001,  
        "max_iterations": 500,  
        "early_stopping": true  
    }  
}
```

**Training Response:**

```
{  
    "status": "success",  
    "metrics": {  
        "r2_score": 0.9872,  
        "mse": 0.0123,  
        "mae": 0.0891,  
        "rmse": 0.1109,  
        "training_time": 23.45  
    },  
    "hyperparameters": {  
        "hidden_layer_1": 128,  
        "hidden_layer_2": 64,  
        "hidden_layer_3": 32,  
        "learning_rate": 0.001,  
        "max_iterations": 500,  
        "early_stopping": true  
    }  
}
```

**Prediction Response:**

```
{  
    "predictions": [1.234, 5.678, ...],
```

(continues on next page)

(continued from previous page)

```

"count": 100
}

```

## 4.11.6 Data Processing Pipeline

### Data Validation

#### Step 1: File Format Validation

```

# Check file extension
if not filename.endswith('.pkl'):
    raise ValueError("File must be .pkl format")

# Try to load pickle
try:
    data = pickle.load(file)
except Exception:
    raise ValueError("Invalid pickle file")

```

#### Step 2: Structure Validation

```

# Training data
if not isinstance(data, dict):
    raise ValueError("Must be dictionary")

if 'X' not in data or 'y' not in data:
    raise ValueError("Must contain 'X' and 'y'")

# Prediction data
if not isinstance(data, np.ndarray):
    raise ValueError("Must be NumPy array")

```

#### Step 3: Shape Validation

```

# Check dimensions
if data['X'].shape[1] != 5:
    raise ValueError("X must have 5 features")

if data['y'].ndim != 1:
    raise ValueError("y must be 1D")

if data['X'].shape[0] != data['y'].shape[0]:
    raise ValueError("X and y must have same samples")

```

#### Step 4: Value Validation

```

# Check for invalid values
if np.any(np.isnan(data['X'])) or np.any(np.isinf(data['X'])):
    raise ValueError("X contains NaN or inf")

if np.any(np.isnan(data['y'])) or np.any(np.isinf(data['y'])):
    raise ValueError("y contains NaN or inf")

```

## Data Preprocessing

### Step 1: Clean Data

```
# Remove NaN rows
mask = ~(np.isnan(X).any(axis=1) | np.isnan(y))
X = X[mask]
y = y[mask]
```

### Step 2: Split Data

```
# 60% train, 20% validation, 20% test
X_temp, X_test, y_temp, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

X_train, X_val, y_train, y_val = train_test_split(
    X_temp, y_temp, test_size=0.25, random_state=42
)
```

### Step 3: Standardize

```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_val = scaler.transform(X_val)
X_test = scaler.transform(X_test)
```

## 4.11.7 Security Considerations

### Current Implementation

#### Suitable for:

- Local development
- Coursework/academic use
- Single-user scenarios

#### Not suitable for:

- Production deployment
- Multi-user systems
- Public internet exposure

### Security Measures

#### Input Validation:

- File size limits
- Format validation
- Value range checking
- Type validation with Pydantic

#### CORS Configuration:

```
app.add_middleware(
    CORSMiddleware,
    allow_origins=["http://localhost:3000"],
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)
```

**For Production:**

Would need:

- Authentication & authorization
- Rate limiting
- File scanning
- HTTPS/TLS
- Database for state
- Session management
- Input sanitization
- Error message sanitization

## 4.11.8 Scalability

### Current Limitations

- Single server instance
- In-memory state
- No horizontal scaling
- No load balancing
- Limited to CPU training

### Potential Improvements

#### For Higher Scale:

##### 1. Database Integration:

- Store models in database
- Persist training state
- Support multiple users

##### 2. Queue System:

- Background job processing
- Async training tasks
- Progress tracking

##### 3. Caching:

- Redis for session state

- Model caching
- Result caching

**4. Microservices:**

- Separate training service
- Separate prediction service
- API gateway

**5. GPU Support:**

- PyTorch/TensorFlow
- CUDA acceleration
- Larger networks

### 4.11.9 Deployment Options

#### Development

**Local:**

```
./scripts/deploy-local.sh
```

**Docker:**

```
docker compose up
```

#### Production

**Cloud Platforms:**

- AWS (ECS, Lambda, SageMaker)
- Google Cloud (Cloud Run, AI Platform)
- Azure (App Service, ML)

**Containerization:**

- Docker images
- Kubernetes orchestration
- Auto-scaling

### 4.11.10 Monitoring & Logging

#### Current Logging

**Backend:**

- Uvicorn access logs
- Python print statements
- Error stack traces

**Frontend:**

- Console.log debugging

- Error boundaries

### Production Logging

Would need:

- Structured logging (JSON)
- Log aggregation (ELK stack)
- Error tracking (Sentry)
- Performance monitoring (APM)
- User analytics

### 4.11.11 Testing Strategy

**Unit Tests:**

- Backend endpoints (pytest)
- Neural network module
- Data handlers
- Validation logic

**Integration Tests:**

- End-to-end workflows
- API contract testing
- Database interactions

**Coverage:**

- 74.54% overall
- 52 total tests

See [Testing Overview](#) for details.

### 4.11.12 Next Steps

- [Local Deployment Guide](#) - Local deployment guide
- deployment/docker - Docker deployment
- [Backend API Reference](#) - Backend API reference
- [Testing Overview](#) - Testing documentation



# 5

## Indices and Tables

- genindex
- modindex
- search



# 6

## Project Information

**Author**

Makimona Kiakisolako (bamk3)

**Institution**

University of Cambridge

**Course**

DIS Course 2025

**License**

MIT

**Version**

0.1.0

**Contact**

bamk3@cam.ac.uk

### 6.1 Links

- *GitHub Repository*
- API Documentation<sup>4</sup>
- fivedreg Package Documentation
- *Issue Tracker*

---

<sup>4</sup> <http://localhost:8000/docs>



## Python Module Index

f

fivedreg.base\_fivedreg, 50  
fivedreg.data\_hand.module, 52